

# Time Series Forecasting Project

## ARIMA & SARIMAX - Workflow



# ARIMA and Seasonal ARIMA

## Autoregressive Integrated Moving Averages

The general process for ARIMA models is the following:

Visualize the Time Series Data

Make the time series data stationary

Plot the Correlation and AutoCorrelation Charts

Construct the ARIMA Model or Seasonal ARIMA based on the data

Use the model to make predictions

Let's go through these steps!

```
In [1]: import numpy as np  
import pandas as pd
```

```
import matplotlib.pyplot as plt  
%matplotlib inline
```

```
In [3]: df=pd.read_csv('perrin-freres-monthly-champagne-.csv')
```

```
In [4]: df.head()
```

```
Out[4]:   Month  Perrin Freres monthly champagne sales millions ?64-?72  
0  1964-01          2815.0  
1  1964-02          2672.0  
2  1964-03          2755.0  
3  1964-04          2721.0  
4  1964-05          2946.0
```

```
In [5]: df.tail()
```

```
Out[5]:   Month  Perrin Freres monthly champagne sales millions ?64-?72  
102           1972-07          4298.0  
103           1972-08          1413.0  
104           1972-09          5877.0  
105             NaN            NaN  
106  Perrin Freres monthly champagne sales millions...      NaN
```

```
In [6]: ## Cleaning up the data  
df.columns=["Month","Sales"]  
df.head()
```

```
Out[6]:   Month  Sales  
0  1964-01  2815.0  
1  1964-02  2672.0  
2  1964-03  2755.0  
3  1964-04  2721.0  
4  1964-05  2946.0
```

```
In [7]: df.tail()
```

```
Out[7]:   Month  Sales  
102           1972-07  4298.0  
103           1972-08  1413.0  
104           1972-09  5877.0  
105             NaN            NaN  
106  Perrin Freres monthly champagne sales millions...      NaN
```

```
In [8]: ## Drop last 2 rows  
df.drop(106, axis=0, inplace=True)
```

```
In [9]: df.tail()
```

```
Out[9]:
```

	Month	Sales
101	1972-06	5312.0
102	1972-07	4298.0
103	1972-08	1413.0
104	1972-09	5877.0
105	NaN	NaN

```
In [11]: df.drop(105, axis=0, inplace=True)
```

```
In [12]: df.tail()
```

```
Out[12]:
```

	Month	Sales
100	1972-05	4618.0
101	1972-06	5312.0
102	1972-07	4298.0
103	1972-08	1413.0
104	1972-09	5877.0

```
In [13]: # Convert Month into Datetime  
df['Month']=pd.to_datetime(df['Month'])
```

```
In [14]: df.head()
```

```
Out[14]:
```

	Month	Sales
0	1964-01-01	2815.0
1	1964-02-01	2672.0
2	1964-03-01	2755.0
3	1964-04-01	2721.0
4	1964-05-01	2946.0

```
In [15]: df.set_index('Month', inplace=True)
```

```
In [16]: df.head()
```

```
Out[16]:
```

	Sales
1964-01-01	2815.0
1964-02-01	2672.0
1964-03-01	2755.0
1964-04-01	2721.0
1964-05-01	2946.0

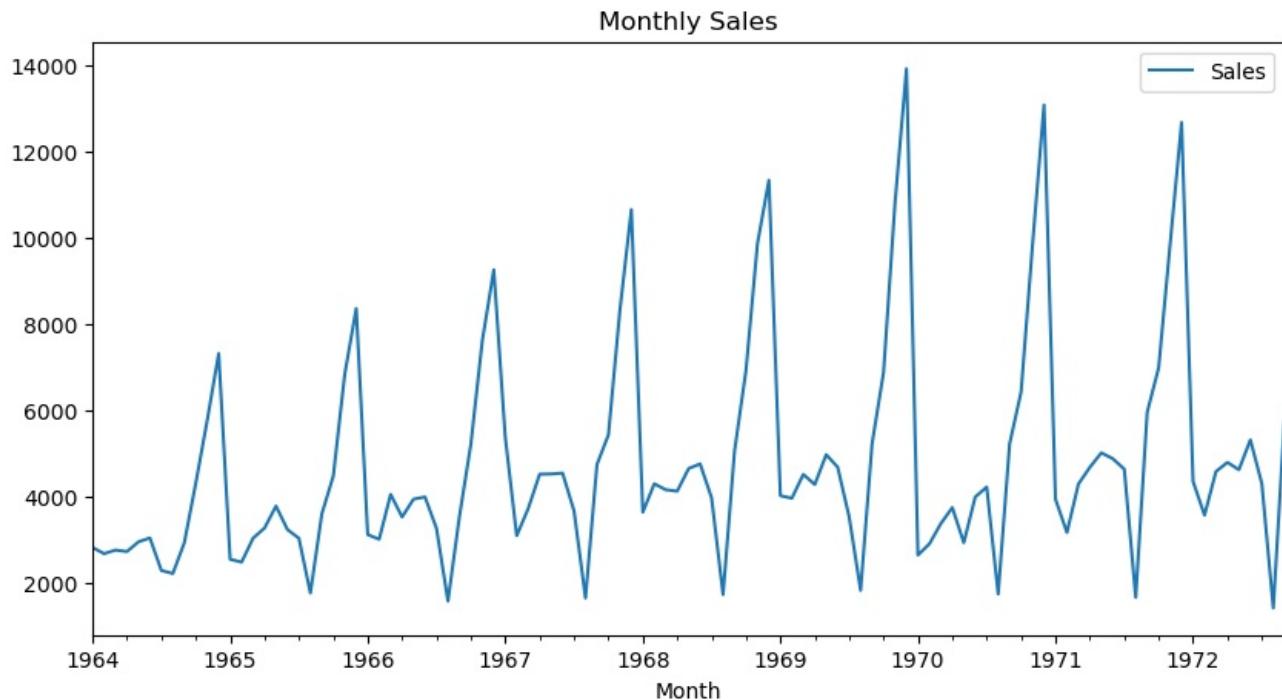
```
In [17]: df.describe()
```

```
Out[17]:
```

	Sales
count	105.000000
mean	4761.152381
std	2553.502601
min	1413.000000
25%	3113.000000
50%	4217.000000
75%	5221.000000
max	13916.000000

## Step 2: Visualize the Data

```
In [23]: df.plot(figsize=(10,5), title="Monthly Sales")
plt.show()
```



```
In [24]: ### Testing For Stationarity
```

```
from statsmodels.tsa.stattools import adfuller
```

```
In [25]: test_result=adfuller(df['Sales'])
```

```
In [26]: #Ho: It is non stationary
#H1: It is stationary
```

```
def adfuller_test(sales):
    result=adfuller(sales)
    labels = ['ADF Test Statistic','p-value','#Lags Used','Number of Observations Used']
    for value,label in zip(result,labels):
        print(label+' : '+str(value) )
    if result[1] <= 0.05:
        print("strong evidence against the null hypothesis(Ho), reject the null hypothesis. Data has no unit root")
    else:
        print("weak evidence against null hypothesis, time series has a unit root, indicating it is non-stationary")
```

```
In [27]: adfuller_test(df['Sales'])
```

```
ADF Test Statistic : -1.833593056327635
p-value : 0.36391577166023914
#Lags Used : 11
Number of Observations Used : 93
weak evidence against null hypothesis, time series has a unit root, indicating it is non-stationary
```

## Differencing

```
In [28]: df['Sales First Difference'] = df['Sales'] - df['Sales'].shift(1)
```

```
In [29]: df['Sales'].shift(1)
```

```
Out[29]: Month
1964-01-01      NaN
1964-02-01    2815.0
1964-03-01    2672.0
1964-04-01    2755.0
1964-05-01    2721.0
...
1972-05-01    4788.0
1972-06-01    4618.0
1972-07-01    5312.0
1972-08-01    4298.0
1972-09-01    1413.0
Name: Sales, Length: 105, dtype: float64
```

```
In [30]: df['Seasonal First Difference']=df['Sales']-df['Sales'].shift(12)
```

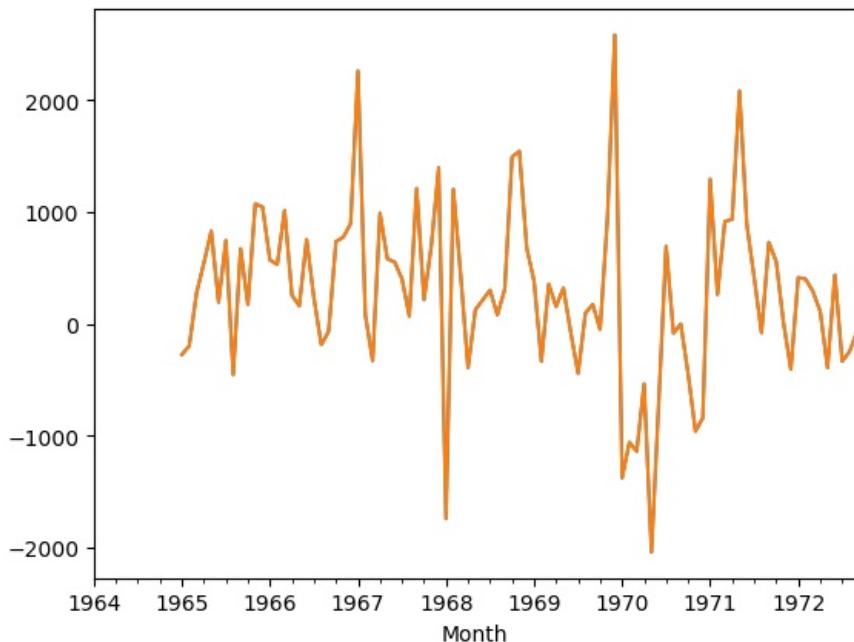
```
In [31]: df.head(15)
```

```
Out[31]:   Sales  Sales First Difference  Seasonal First Difference
Month
1964-01-01  2815.0           NaN           NaN
1964-02-01  2672.0        -143.0           NaN
1964-03-01  2755.0          83.0           NaN
1964-04-01  2721.0         -34.0           NaN
1964-05-01  2946.0         225.0           NaN
1964-06-01  3036.0          90.0           NaN
1964-07-01  2282.0        -754.0           NaN
1964-08-01  2212.0         -70.0           NaN
1964-09-01  2922.0         710.0           NaN
1964-10-01  4301.0        1379.0           NaN
1964-11-01  5764.0        1463.0           NaN
1964-12-01  7312.0        1548.0           NaN
1965-01-01  2541.0        -4771.0        -274.0
1965-02-01  2475.0         -66.0        -197.0
1965-03-01  3031.0         556.0          276.0
```

```
In [32]: ## Again test dickey fuller test
adfuller_test(df['Seasonal First Difference'].dropna())
```

```
ADF Test Statistic : -7.626619157213163
p-value : 2.060579696813685e-11
#Lags Used : 0
Number of Observations Used : 92
strong evidence against the null hypothesis(Ho), reject the null hypothesis. Data has no unit root and is stationary
```

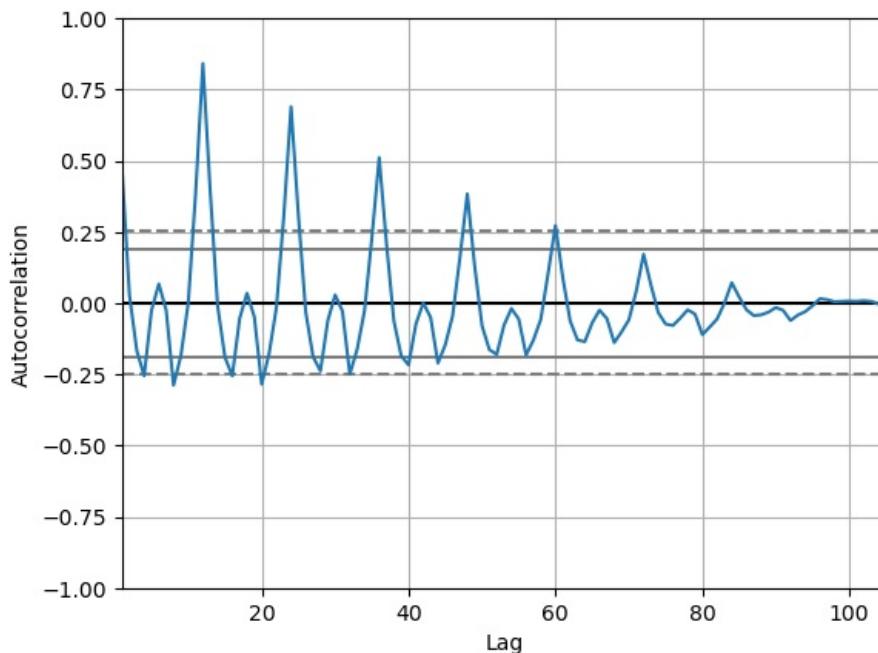
```
In [34]: import matplotlib.pyplot as plt
df['Seasonal First Difference'].plot()
plt.show()
```



## Auto Regressive Model

```
In [37]: from pandas.plotting import autocorrelation_plot
import matplotlib.pyplot as plt

autocorrelation_plot(df['Sales'])
plt.show()
```



This means:

- ✓ data has seasonality
- ✓ Past values affect future values

## Final Thoughts on Autocorrelation and Partial Autocorrelation

Identification of an AR model is often best done with the PACF.

For an AR model, the theoretical PACF “shuts off” past the order of the model. The phrase “shuts off” means that in theory the partial autocorrelations are equal to 0 beyond that point. Put another way, the number of non-zero partial autocorrelations gives the order of the AR model. By the “order of the model” we mean the most extreme lag of  $x$  that is used as a predictor.

Identification of an MA model is often best done with the ACF rather than the PACF.

For an MA model, the theoretical PACF does not shut off, but instead tapers toward 0 in some manner. A clearer pattern for an MA model is in the ACF. The ACF will have non-zero autocorrelations only at lags involved in the model.

p,d,q p AR model lags d differencing q MA lags

```
In [38]: from statsmodels.graphics.tsaplots import plot_acf,plot_pacf
```

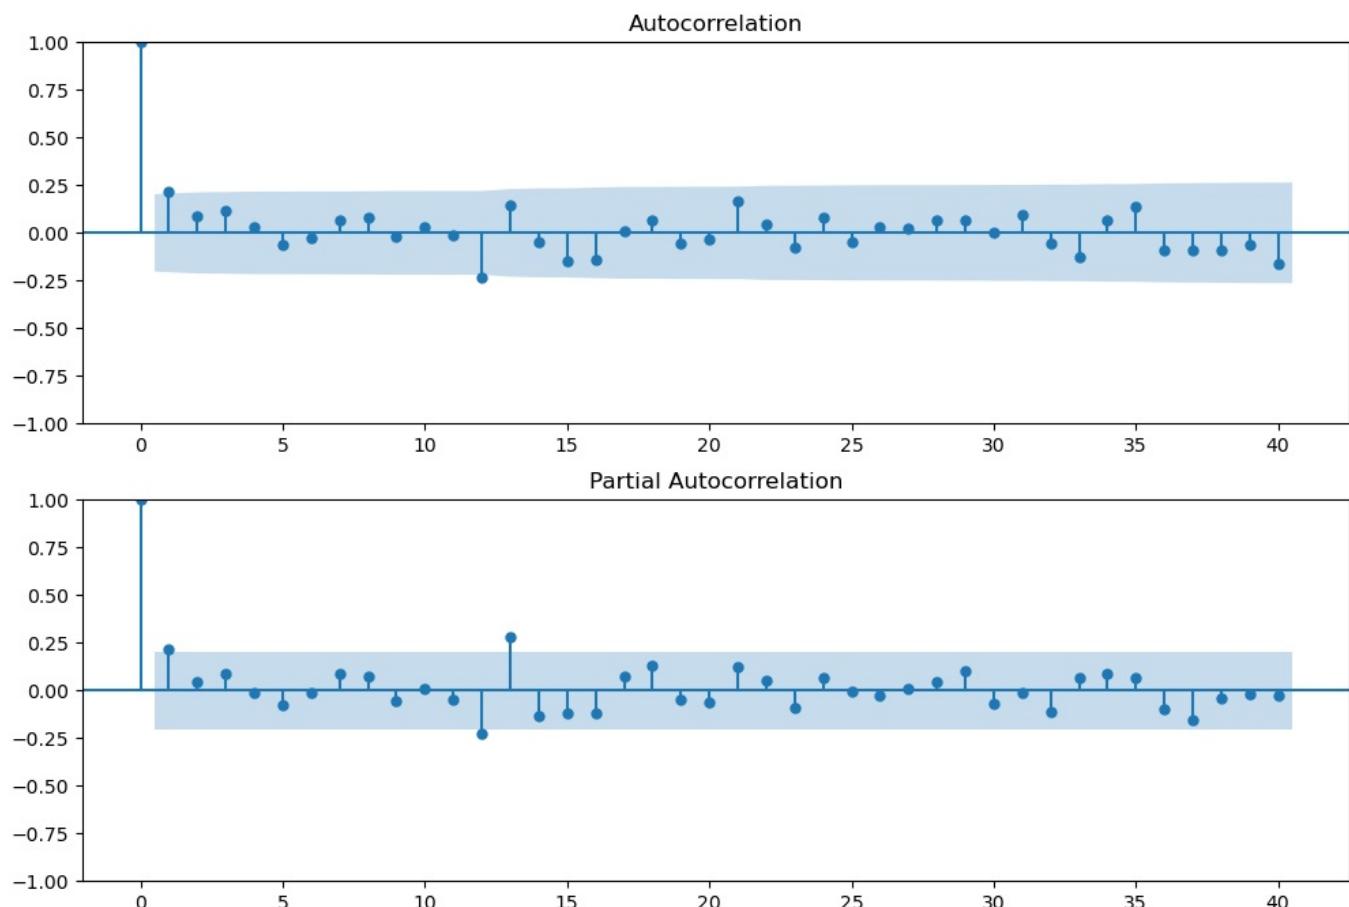
```
In [42]: import matplotlib.pyplot as plt
import statsmodels.api as sm
```

```
fig = plt.figure(figsize=(12,8))

ax1 = fig.add_subplot(211)
sm.graphics.tsa.plot_acf(
    df['Seasonal First Difference'].dropna(),
    lags=40,
    ax=ax1
)

ax2 = fig.add_subplot(212)
sm.graphics.tsa.plot_pacf(
    df['Seasonal First Difference'].dropna(),
    lags=40,
    ax=ax2
)

plt.show()
```



```
In [53]: # For non-seasonal data
#p=1, d=1, q=0 or 1
from statsmodels.tsa.arima_model import ARIMA
```

```
In [54]: df.index = pd.to_datetime(df.index)
```

```
In [55]: df = df.asfreq('MS')
```

```
In [56]: from statsmodels.tsa.arima.model import ARIMA

model = ARIMA(df['Sales'], order=(1,1,1))
model_fit = model.fit()
```

```
In [57]: model_fit.summary()
```

Out[57]:

## SARIMAX Results

<b>Dep. Variable:</b>	Sales	<b>No. Observations:</b>	105			
<b>Model:</b>	ARIMA(1, 1, 1)	<b>Log Likelihood</b>	-952.814			
<b>Date:</b>	Thu, 18 Dec 2025	<b>AIC</b>	1911.627			
<b>Time:</b>	21:23:55	<b>BIC</b>	1919.560			
<b>Sample:</b>	01-01-1964	<b>HQIC</b>	1914.841			
	- 09-01-1972					
<b>Covariance Type:</b>	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.4545	0.114	3.998	0.000	0.232	0.677
ma.L1	-0.9666	0.056	-17.305	0.000	-1.076	-0.857
sigma2	5.226e+06	6.17e+05	8.473	0.000	4.02e+06	6.43e+06
Ljung-Box (L1) (Q):	0.91	Jarque-Bera (JB):	2.59			
Prob(Q):	0.34	Prob(JB):	0.27			
Heteroskedasticity (H):	3.40	Skew:	0.05			
Prob(H) (two-sided):	0.00	Kurtosis:	3.77			

Warnings:

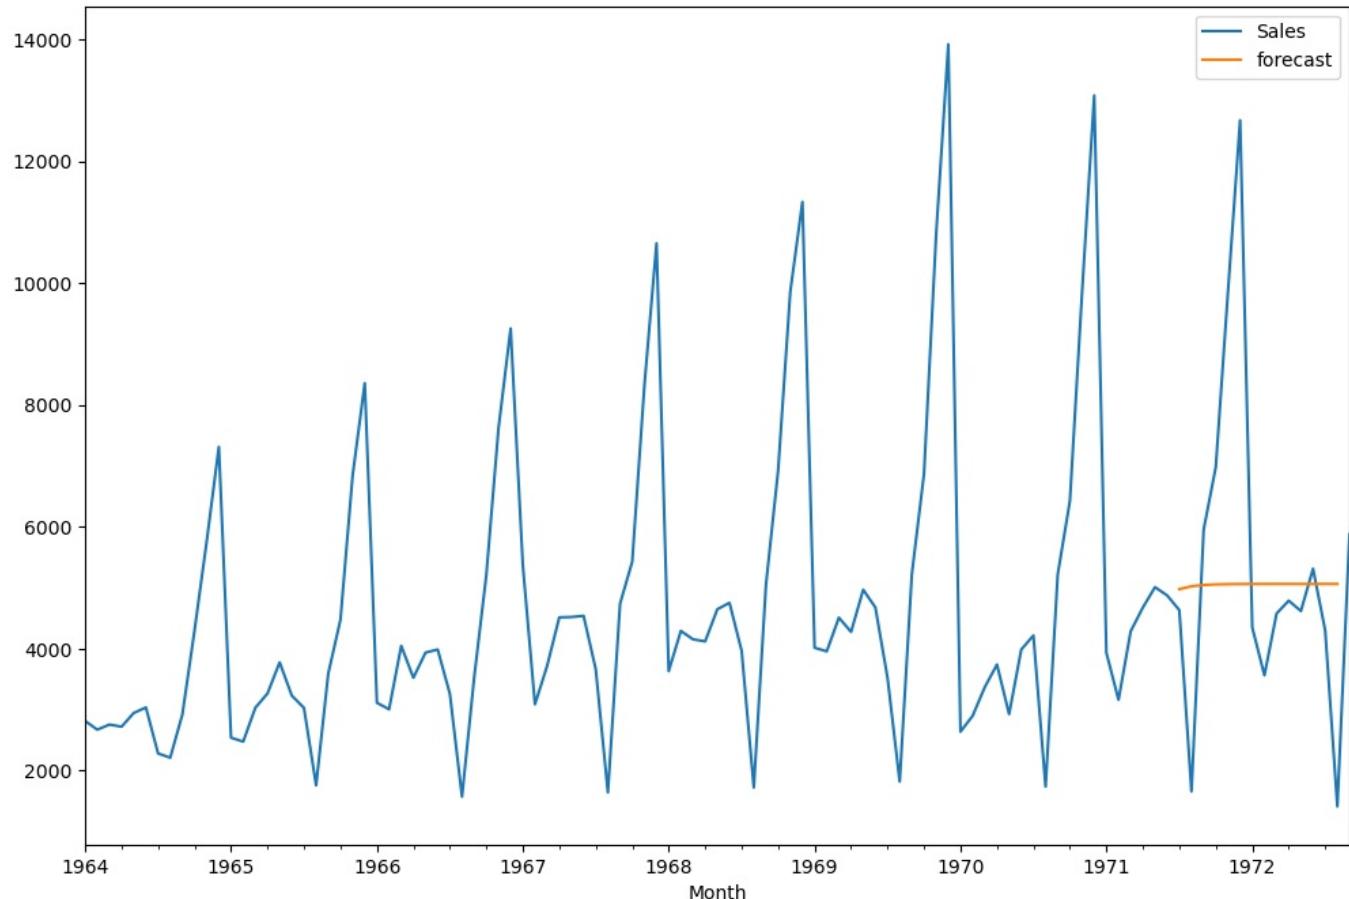
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

In [60]:

```
import matplotlib.pyplot as plt

df['forecast'] = model_fit.predict(start=90, end=103, dynamic=True)

df[['Sales', 'forecast']].plot(figsize=(12,8))
plt.show()
```



In [61]:

```
import statsmodels.api as sm
```

In [62]:

```
model=sm.tsa.statespace.SARIMAX(df['Sales'],order=(1, 1, 1),seasonal_order=(1,1,1,12))
results=model.fit()
```

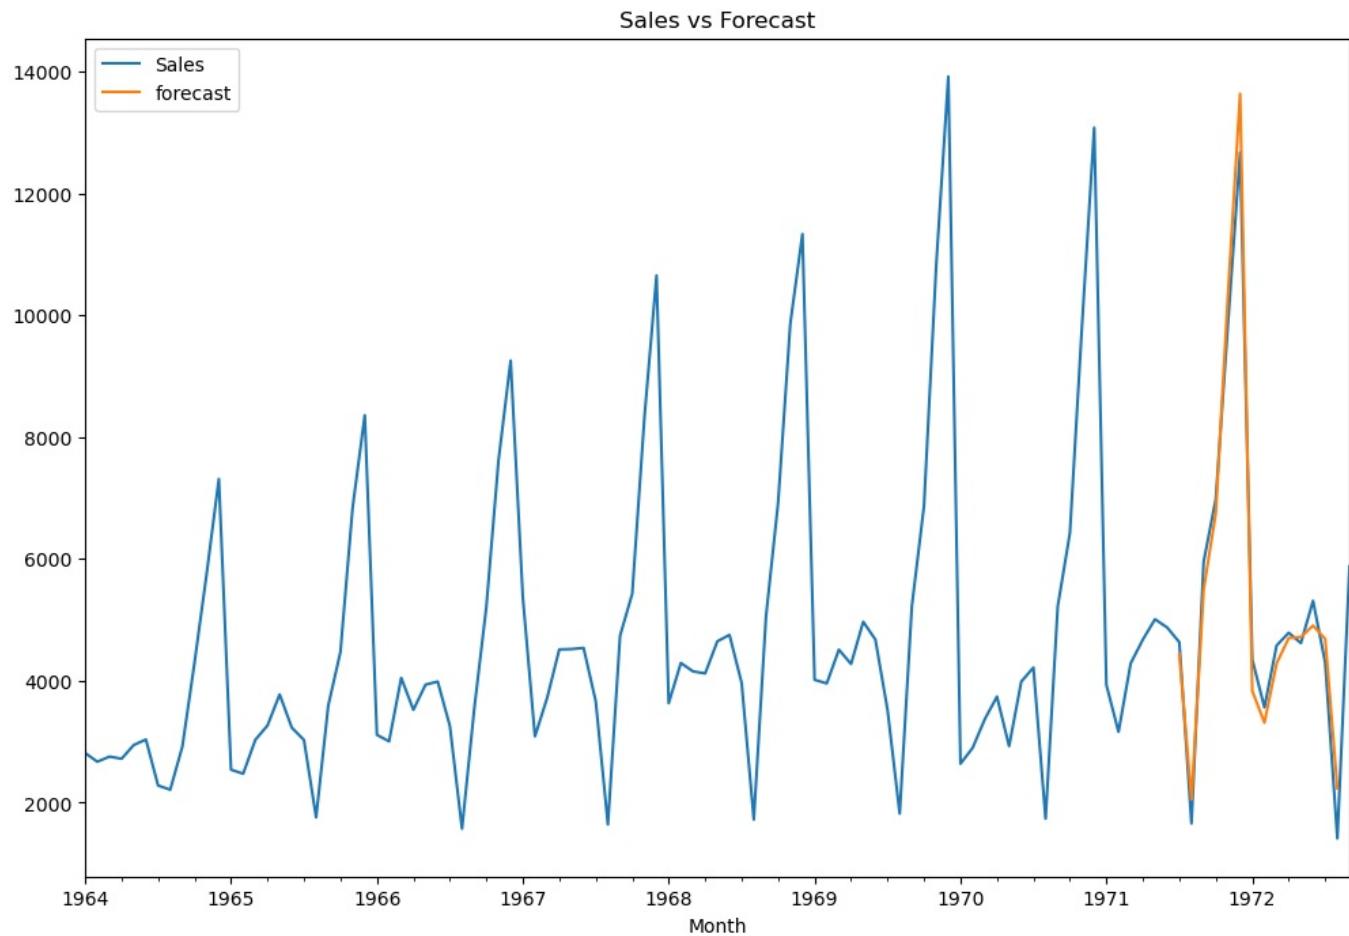
```
In [64]: %matplotlib inline
```

```
In [66]: import matplotlib.pyplot as plt

df['forecast'] = results.predict(start=90, end=103, dynamic=True)

ax = df[['Sales', 'forecast']].plot(figsize=(12,8))
ax.set_title("Sales vs Forecast")

plt.show()
```



```
In [67]: from pandas.tseries.offsets import DateOffset
future_dates=[df.index[-1]+DateOffset(months=x) for x in range(0,24)]
```

```
In [68]: future_datest_df=pd.DataFrame(index=future_dates[1:],columns=df.columns)
```

```
In [69]: future_datest_df.tail()
```

```
Out[69]:
```

	Sales	Sales First Difference	Seasonal First Difference	forecast
1974-04-01	NaN	NaN	NaN	NaN
1974-05-01	NaN	NaN	NaN	NaN
1974-06-01	NaN	NaN	NaN	NaN
1974-07-01	NaN	NaN	NaN	NaN
1974-08-01	NaN	NaN	NaN	NaN

```
In [75]: future_df=pd.concat([df,future_datest_df])
```

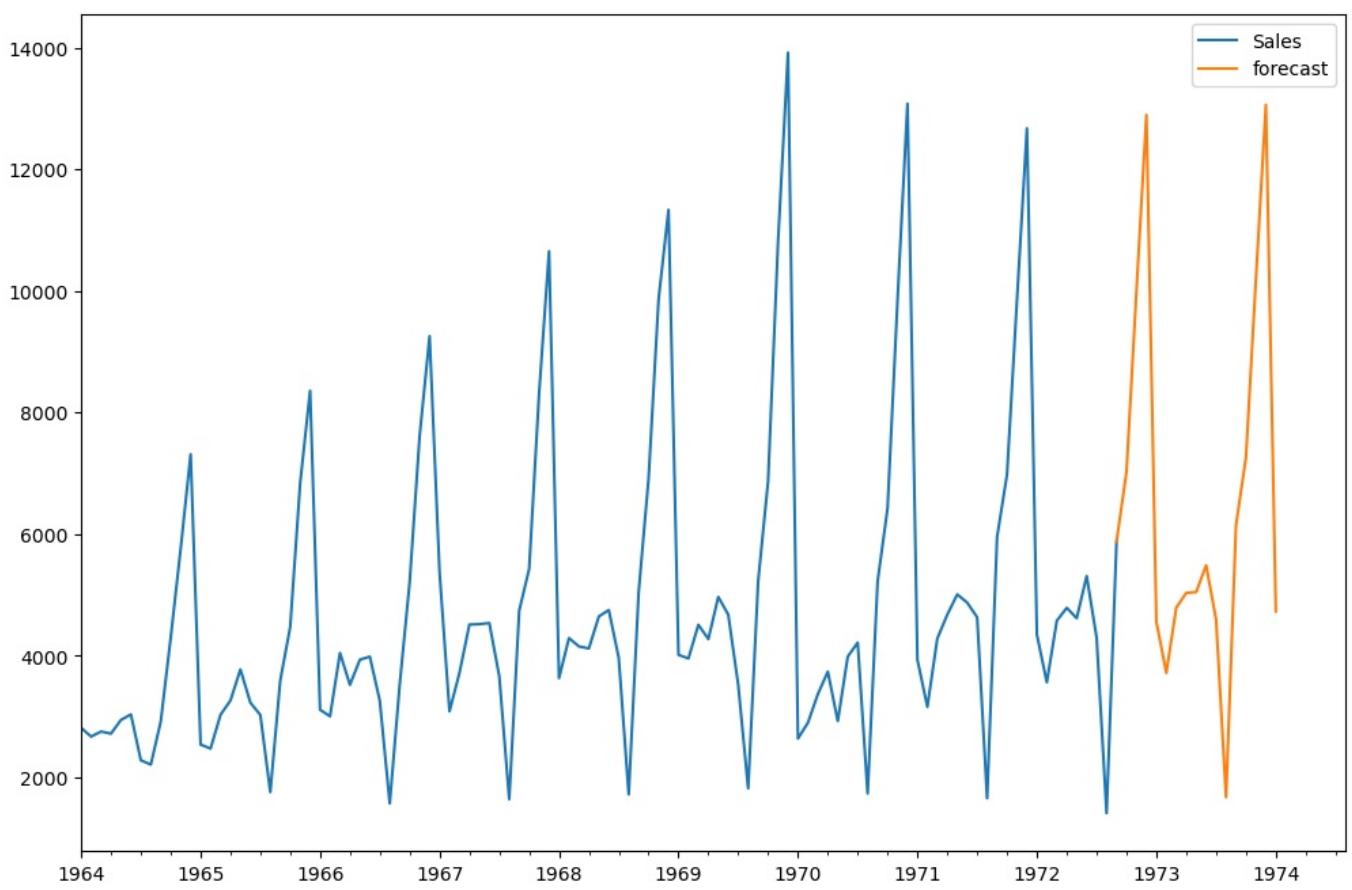
C:\Users\asus1\AppData\Local\Temp\ipykernel\_17052\1723908026.py:1: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the concat operation.

```
future_df=pd.concat([df,future_datest_df])
```

```
In [76]: import matplotlib.pyplot as plt

future_df['forecast'] = results.predict(start=104, end=120, dynamic=True)

future_df[['Sales','forecast']].plot(figsize=(12,8))
plt.show()
```



## Top 5 Project Insights

- ① Sales data shows strong seasonality and trend, making time series forecasting models like ARIMA and SARIMAX suitable for prediction.
- ② Stationarity checking and differencing were crucial, as ARIMA models perform best when the time series has stable statistical properties.
- ③ ACF and PACF analysis helped identify optimal AR and MA parameters, improving model structure and prediction accuracy.
- ④ SARIMAX performed better than basic ARIMA by capturing seasonal patterns present in the sales data.
- ⑤ Forecast results followed the overall sales trend, proving that time series models are effective baseline approaches for future sales planning.

## Project Conclusion

In this project, time series forecasting was performed to predict future sales using ARIMA and SARIMAX models.

The data was analyzed for trends, seasonality, and stationarity, followed by proper preprocessing and model selection.

ARIMA successfully captured the overall sales pattern, while SARIMAX provided improved performance by handling seasonality.

The results demonstrate that time series models are effective for short-term forecasting based on historical data.

Future accuracy can be enhanced by incorporating external factors and more advanced forecasting techniques.

The data showed clear seasonality and trend. Proper preprocessing and parameter selection improved forecasting, and SARIMAX provided better results by handling seasonality compared to ARIMA

In [ ]:

Loading [MathJax]/extensions/Safe.js



THE END

---

# THANK YOU!

We appreciate your time and attention!

